


ATHS – Post-Preprocessing Development Guide

Prompt Injection Threat Hypothesis System
Threat Detection → Hypothesis Generation → Threat Scoring → Dashboard/XAI

Purpose: This guide gives the team a practical sequence for continuing ATHS after the completed Data Preprocessing Layer.

Current preprocessing outputs: dataset/processed/features.csv and dataset/processed/embeddings.npy (4391 × 384 in the current dataset).

1. Overall ATHS Architecture



2. Clear Steps After the Preprocessing Layer

- **Step 1:** Load features.csv and embeddings.npy.
- **Step 2:** Create train/validation/test splits. Prefer group-aware splitting using group_id to reduce leakage.
- **Step 3:** Train a simple baseline classifier.
- **Step 4:** Train a stronger non-linear classifier and compare results.
- **Step 5:** Add the 384-D semantic embeddings and compare performance again.
- **Step 6:** Add the interpretable security-rule layer.
- **Step 7:** Fuse ML probability, rule evidence and semantic evidence into a calibrated threat score.
- **Step 8:** Generate and rank attack hypotheses.
- **Step 9:** Define one structured JSON result for the dashboard.
- **Step 10:** Build the Streamlit/XAI dashboard and connect the real engine.

3. Threat Detection Engine – Recommended Models

Model / method	Why use it
Logistic Regression	Fast, interpretable baseline; gives probability output and shows whether engineered features already contain useful s
Random Forest	Handles non-linear interactions among binary/numeric security features and provides feature importance.
XGBoost / LightGBM (optional)	Strong tabular-data comparison model if permitted; use it only after establishing simpler baselines.
384-D embeddings + cosine similarity	Captures semantic similarity and helps detect rephrased injections with different wording.
Rule-based detector	Turns preprocessing indicators into transparent security evidence.
SHAP	Explains supported ML predictions by showing feature contributions.

4. Why Use a Hybrid Detector?

A single classifier only answers whether a prompt looks malicious. ATHS should combine predictive and explainable signals. Rules provide explicit evidence; ML learns patterns from the labelled dataset; embeddings help with semantic/rephrasing variation.

```

Handcrafted/security features
+
384-D embedding
↓
ML prediction
+
Rule evidence
+
Semantic similarity
↓
Final threat score

```

5. Training and Evaluation

- Train Logistic Regression as the baseline.
- Train Random Forest and compare it with the baseline.
- Optionally evaluate XGBoost/LightGBM.
- Compare handcrafted-only features against handcrafted + embeddings.
- Tune model/decision thresholds on the validation set.
- Keep the test set untouched until final evaluation.

6. Metrics to Report

Accuracy, precision, recall, F1-score, ROC-AUC, confusion matrix, false-positive rate and false-negative rate. Security evaluation should pay particular attention to false negatives because they represent attacks that were missed.

7. Threat Scoring

Use a transparent fusion score initially, then tune/calibrate the contributions and thresholds using validation data. Do not present arbitrary weights as scientifically validated.

Conceptual fusion:

```

Threat Score =
  calibrated ML probability contribution
+ rule-evidence contribution
+ semantic-similarity contribution

```

Normalize to 0-100.

Example starting interpretation:

```

0-29   LOW
30-59  MEDIUM
60-79  HIGH
80-100 CRITICAL

```

Tune these thresholds with validation data.

Example evidence: instruction_override, system_prompt_reference, data_extraction_attempt, role_change_attempt, tool_manipulation, policy_bypass, prompt_leakage, sensitive_data_reference and obfuscation.

8. Hypothesis Generation Engine

The Threat Detection Engine answers “**Is this suspicious?**”. The Hypothesis Engine answers “**What attack could explain the evidence, and how confident are we?**”.

```

Detection output
↓
Generate candidate hypotheses
↓
Match evidence to each hypothesis
↓
Calculate confidence
↓
Rank hypotheses
↓
Return hypotheses + evidence + confidence

```

9. Suggested Hypotheses

Hypothesis	Useful evidence
Direct Prompt Injection	instruction_override, imperative_signal, semantic similarity
System Prompt Extraction	system_prompt_reference, data_extraction_attempt, prompt_leakage
Role Manipulation	role_change_attempt, instruction_override
Tool Manipulation	tool_manipulation, data_extraction_attempt
Obfuscated Injection	obfuscation_score + injection indicators
Indirect Prompt Injection	injection indicators inside untrusted retrieved/tool content
RAG / Context Poisoning	suspicious instructions in retrieved/external context
Multi-turn / Escalation	suspicious progression across conversation history

10. Example Hypothesis Output

```
{
  "threat_detected": true,
  "threat_score": 91,
  "severity": "HIGH",
  "hypotheses": [
    {
      "type": "direct_prompt_injection",
      "confidence": 0.94,
      "evidence": [
        "instruction_override",
        "system_prompt_reference",
        "data_extraction_attempt"
      ]
    }
  ],
  "recommended_action": "BLOCK"
}
```

11. Context-Aware Detection

Use the trusted/untrusted context produced by preprocessing. A user prompt can be harmless while a retrieved document or tool output contains an injection. The hypothesis engine should therefore record both the attack type and the suspicious source.

Example:

```
Hypothesis: Indirect Prompt Injection
Source: UNTRUSTED retrieved document
Evidence:
- instruction override
- system prompt reference
Confidence: HIGH
```

12. Dashboard/XAI Technology Stack

Technology	Role
Streamlit	Browser-based web dashboard built directly in Python; recommended first implementation.
Plotly	Interactive charts for threat scores, severity, categories, feature contributions and model performance.
SHAP	Feature-level model explanations showing which inputs pushed a prediction toward malicious/benign.
Pandas	Dataset and analytics handling.
NumPy	Embedding and numeric array handling.
scikit-learn	ML models, metrics, confusion matrix and evaluation.
FastAPI + React (optional later)	Separate backend/frontend architecture if the project later needs a production-style web application.

13. Dashboard Pages

- 1. Overview** — Total prompts, threats detected, benign/malicious distribution, high/critical counts, attack categories and severity.
- 2. Live Prompt Analysis** — Prompt input, Analyze button, threat status, score, severity, top hypothesis and recommended action.
- 3. XAI / Explainability** — SHAP/feature contributions, triggered rules, semantic similarity and supporting evidence.
- 4. Threat Hypotheses** — Ranked hypotheses, confidence, evidence, severity and explanation for each.
- 5. Context Analysis** — Trusted vs untrusted context, suspicious retrieved documents/tool outputs and attack source.
- 6. Dataset / Analytics** — Category, severity, obfuscation and feature distributions.
- 7. Model Performance** — Precision, recall, F1, ROC-AUC, confusion matrix and model comparison.

14. What the XAI Layer Must Explain

- **What happened?** Benign, suspicious or malicious.
- **How dangerous?** Threat score and severity.
- **What attack?** Ranked hypotheses.
- **Why?** Security indicators and model contributions.
- **What evidence?** Triggered rules and relevant context.
- **Where did it originate?** User prompt, retrieved document, tool output or conversation.
- **What should happen?** Allow, review or block.

15. Recommended Development Order

```
PHASE 1 → Train/evaluate ML baselines
PHASE 2 → Add Random Forest / optional boosting
PHASE 3 → Add semantic embedding signal
PHASE 4 → Build rule detector
PHASE 5 → Fuse ML + rules + semantic evidence
PHASE 6 → Build hypothesis knowledge base
PHASE 7 → Generate/rank hypotheses
PHASE 8 → Define one JSON result contract
PHASE 9 → Build Streamlit dashboard
PHASE 10 → Add Plotly + SHAP + context visualization
PHASE 11 → Evaluate the complete end-to-end ATHS
```

16. Recommended Folder Structure


```
ATHS/
├── preprocessing/
├── detection/
│   ├── feature_loader.py
│   ├── rule_detector.py
│   ├── ml_detector.py
│   ├── embedding_detector.py
│   └── threat_fusion.py
├── hypothesis/
│   ├── hypothesis_engine.py
│   ├── hypothesis_rules.json
│   ├── evidence_analyzer.py
│   └── hypothesis_ranker.py
├── models/
├── dataset/processed/
│   ├── features.csv
│   └── embeddings.npy
├── evaluation/
└── dashboard/
    ├── app.py
    ├── pages/
    └── xai/
```

17. Dashboard-to-Engine Interface

Agree on a single structured result before integrating the dashboard. The dashboard should consume the result and visualize it, rather than depending on internal model implementation details.

```
{
  "input": {"prompt": "..."},
  "detection": {
    "label": "malicious",
    "score": 0.91,
    "severity": "high"
  },
  "hypotheses": [
    {
      "name": "direct_prompt_injection",
      "confidence": 0.94,
      "evidence": ["instruction_override"]
    }
  ],
  "explainability": {
    "feature_contributions": {},
    "semantic_similarity": [],
    "rules_triggered": []
  },
  "context": {
    "source": "user_prompt",
    "trust_level": "untrusted"
  },
  "recommendation": {"action": "block"}
}
```

18. End-to-End Example



19. Team Deliverables

Team member/layer	Deliverable
Preprocessing	features.csv, embeddings.npy, context/trust representation and feature documentation.
Threat Detection	Trained models, saved model, evaluation metrics and threat probability output.
Hypothesis Engine	Hypothesis rules/knowledge base, candidate generation, confidence/ranking and evidence output.
Dashboard/XAI	Streamlit pages, Plotly visualizations, SHAP explanations, context visualization and final decision display.
Integration	One ATHS analysis interface/API contract and end-to-end test cases.

20. Important Implementation Notes

- Do not display all 384 embedding values in the dashboard. Use embeddings for models/similarity; show human-readable semantic matches instead.
- Do not claim manually chosen threat-score weights are validated. Tune/calibrate them on validation data and document the method.
- Keep the final test set separate until evaluation.
- Use group-aware splitting where related prompts could otherwise appear in both train and test.
- Make explanations faithful to the actual model output and input features.
- The dashboard should visualize engine-generated evidence; it should not invent evidence.
- Start with Streamlit + Plotly + SHAP. A React/FastAPI architecture or browser extension can be added later if time permits.

21. Final Target

The completed ATHS should not merely say “**malicious = 1.**” It should detect the threat, estimate severity, generate competing attack hypotheses, rank them using evidence, explain the decision, identify the source/context of suspicious content, and present the result through a usable browser-based XAI dashboard.